


Which pathology evaluations matter, and how much is enough?

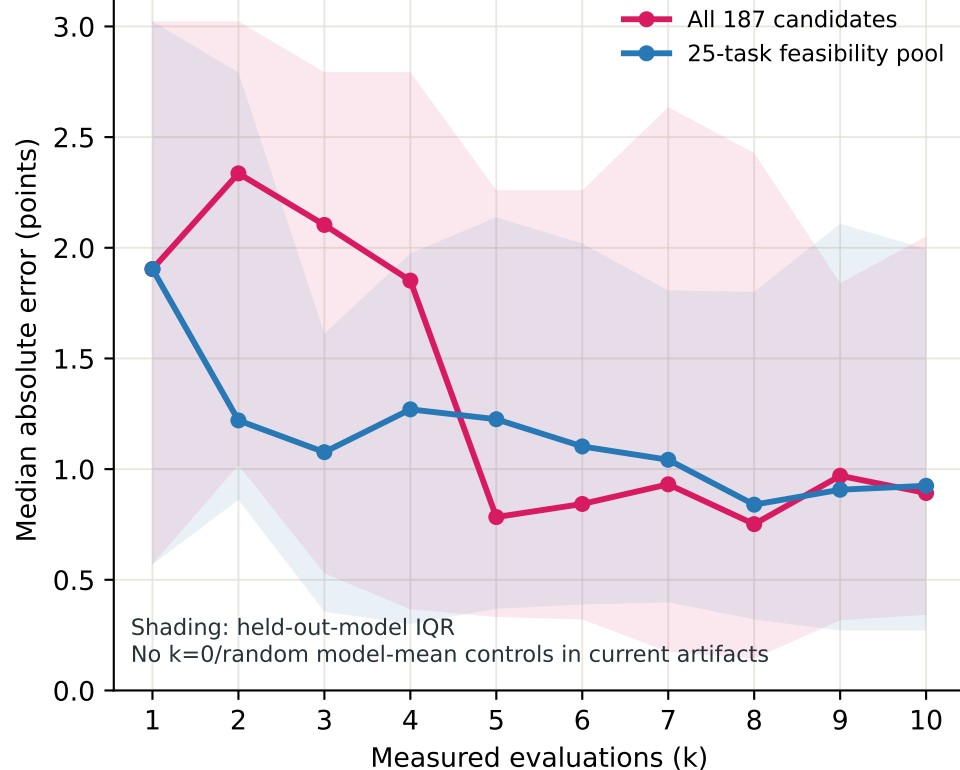
A Retrospective single-task utility

Transductive; exact probe cells included; all 8: THUNDER, 32/59 models each (54%)



B Error predicting mean reported score

Nested prefixes: 41 selection models → 18 held-out models



A ranks retrospective utility—not causal task importance. B predicts each held-out model's mean over its reported normalized scores.
The 25-task pool is a feasibility proxy—not measured cost.