



Which pathology evaluations matter, and how much is enough?

A Retrospective single-task utility

B Error predicting mean reported score



A ranks retrospective utility—not causal task importance. B predicts each held-out model's mean over its reported normalized scores.
The 25-task pool is a feasibility proxy—not measured cost.