


Which pathology evaluations matter, and how much is enough?

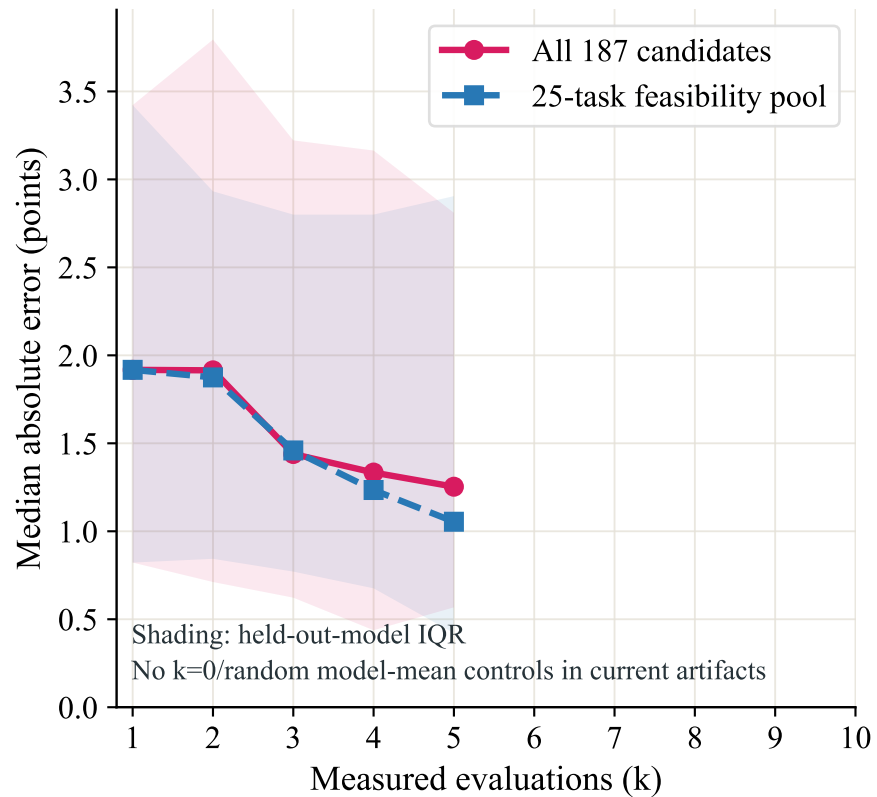
A Retrospective single-task utility

Transductive; exact probe cells included;
all 8: THUNDER, 32/59 models each (54%)



B Error predicting mean reported score

Nested prefixes, leave-one-family-out over 34 family folds;
59 held-out model instances (median 1 validation model/fold)



A ranks retrospective utility—not causal task importance.

B predicts each held-out model's mean over its reported normalized scores.

The 25-task pool is a feasibility proxy—not measured cost.